

AI Code Reviewer Championship

Round Two - four premium reviewers, one frozen production commerce codebase, blind source-verified scoring.

15-17 September 2026 | CR-BENCH-02 | Executive decision brief

CHAMPION - HIGHEST BLIND TECHNICAL SCORE

95.6

ASTRA A+

Frozen technical grade: A. The A+ is the human-facing presentation label.

Bottom line

Astra won on validity, evidence discipline, and unique state/correctness discoveries. Fable delivered the broadest elite review. GLM 5.3 remained strong and actionable. GLM Flash found the architecture, but lost ground on a false concurrency claim and several important misses.

362/362

TARGET INTEGRITY

Every file verified at every checkpoint.

32

ISSUE FAMILIES

Consolidated by the blind referee.

12

SCORING DIMENSIONS

Every model received a full capability fingerprint.

Final leaderboard

01	Astra Frozen technical grade A		95.6	A+
02	Fable 5.1 Blind-scored senior review		90.0	A
03	GLM 5.3 Blind-scored senior review		85.0	B
04	GLM 5.3 Flash Blind-scored senior review		71.2	C

What we actually tested

A natural senior-review benchmark - not a checklist-following contest. Every model received the same frozen target and raw prompt. The rubric, expected findings, prior reviews, and other contestants' work were withheld.

The specimen

DockBloxx production storefront at commit **059ccdba...** 362 regular files, no nested Git, no symlinks, manifest verified 362/362.

The track

Same target. Same prompt. Read-only review. No coaching. No scoring criteria. Raw reports frozen exactly as delivered.

The referee

Sol High scored anonymous A/B/C/D with read-only source access in a physically isolated workspace.

The closeout

Jarvis audited blind. Tony revealed identities only after scoring and adjudication were frozen.

The blind race



Why the contestants did not see the rubric

We wanted natural reviewer judgment: what they chose to inspect, what they treated as material, and how they calibrated uncertainty. Showing the 12 dimensions would have turned the race into rubric optimization.

The 100-point instrument



Four strong reviews. One clear winner.

01	Astra Frozen technical grade A		95.6	A+
02	Fable 5.1 Blind-scored senior review		90.0	A
03	GLM 5.3 Blind-scored senior review		85.0	B
04	GLM 5.3 Flash Blind-scored senior review		71.2	C

95.6

TOP BLIND SCORE

Astra

No material false finding found.

What every reviewer got right

All four identified the same central failure: **the browser is treated as the financial ledger**. Order status, charge amount, discounts, and shipping authority live on the wrong side of the trust boundary.

What actually separated the field

Validity, evidence discipline, calibration, state/concurrency reasoning, and ordinary-user correctness. Architecture did not separate them - all four scored 5/5.

The scoring lesson

One false claim can be expensive

A demonstrably wrong limiter claim pulled down three dimensions at once: validity, confidence, and state/concurrency.

Depth beat breadth at the top

Fable covered more. Astra verified more deeply and found ordinary shopping defects the broader sweep missed.

Source access changed the result

The blind referee could settle disputed claims directly against the frozen target instead of grading plausibility.

Blind adjudication result

Jarvis accepted Sol's scores as written. No score, grade, or rank changed after the audit - and identities were still sealed.

The 12-dimension heatmap

Raw 0-5 referee ratings after reveal. Green = principal-level, blue = strong, amber = competent with notable gaps.

DIMENSION	WT	ASTRA	FABLE	GLM 5.3	FLASH
Finding validity	15	5	4	4	3
Important-issue coverage	15	4	4	4	3
Evidence quality	10	5	5	5	4
Severity calibration	7	4	4	3	3
Confidence / epistemic discipline	7	5	5	4	3
Security / trust-boundary reasoning	8	5	5	5	4
Logic / state / concurrency	8	5	4	4	3
Architecture / root-cause reasoning	8	5	5	5	5
Test-quality reasoning	6	5	5	4	4
Scope discipline / signal-to-noise	5	5	4	4	4
Operational usefulness	6	5	5	5	4
Independent reviewer judgment	5	5	5	4	4
Weighted total	100	95.60	90.00	85.00	71.20

The universal strength

Architecture / root-cause reasoning: 5/5 for all four. Every contestant could diagnose the structural problem. Architecture alone was not enough to win.

The winning edge

Astra's clean sweep across validity, evidence, epistemic discipline, state/concurrency, signal/noise, and judgment produced the field's only near-perfect profile.

Astra vs Fable 5.1

RANK 1 · CHAMPION

Astra**95.6** A+

Validity		5
Coverage		4
Evidence		5
Calibration		4
Epistemic		5
Security		5
State		5
Architecture		5
Tests		5
Signal		5
Usefulness		5
Judgment		5

- **Unique wins:** zero-price complex variations; both pagination defects; exact process-exit trigger; Sharp/Node engine mismatch.
- **Why it won:** no material false claim, best evidence, best state/concurrency work, strongest targeted verification.
- **Meaningful misses:** coupon expiry/dead math, limiter architecture, customer-data persistence, CMS hardening, environment wiring.

RANK 2 · ELITE BREADTH

Fable 5.1**90.0** A

Validity		4
Coverage		4
Evidence		5
Calibration		4
Epistemic		5
Security		5
State		4
Architecture		5
Tests		5
Signal		4
Usefulness		5
Judgment		5

- **Unique wins:** silent account creation/signup analytics; caller-controlled Stripe customer and receipt behavior.
- **Why it excelled:** broad financial/security sweep, precise evidence, excellent test critique, disciplined conditional labels.
- **Meaningful misses:** variation pricing, pagination races, live coupon eligibility math, exact process-exit trigger.

Top-tier contrast

Fable was broader. Astra was cleaner and deeper. The 5.6-point gap came from validity under scrutiny and high-value ordinary-user correctness discoveries, not from architecture or security awareness.

Full model vs fast variant

RANK 3 · STRONG CORE REVIEW

GLM 5.3**85.0_B**

- **Strongest:** evidence, security/trust, architecture, operational usefulness.
- **What held it back:** severity inflation on latent/dead paths, overstated image-secret and blog conclusions.
- **Important misses:** zero-price variation, pagination races, exact termination path, Sharp/Node contract.

RANK 4 · FAST ARCHITECTURE PASS

GLM 5.3 Flash**71.2_C**

- **Strongest:** architecture/root cause at 5/5, clear core checkout/security story.
- **Critical error:** called an active limiter dead code; four call sites prove the opposite.
- **Important misses:** coupon disclosure, process termination, zero-price variation, pagination defects, runtime contract.

The full model bought more quality

GLM 5.3 scored **13.8 points higher** and produced the stronger verified review on this specimen.

Flash still found the architecture

The fast variant was not confused about the big system failure. It lost on verification, omissions, and one confident false claim.

What the reviewers were protecting

32

Consolidated issue families

All four found the central failure: the browser acts as the financial ledger. The referee then separated repeated symptoms, conditional effects, false claims, and unique discoveries.

PAYMENT AUTHORITY

Unauthenticated order mutation from caller-controlled input.

STRIPE AMOUNT

Charge amount and currency accepted without order reconciliation.

DISCOUNT + SHIPPING

Client-supplied coupon metadata, fees, and shipping totals trusted.

NO RECONCILIATION

Browser callback is the supplied paid-order transition; no signed webhook.

PII EXPOSURE

Customer and coupon routes expose more data than the client needs.

CREDENTIAL LOGGING

Commerce API credentials and customer order data reach logs.

PROCESS TERMINATION

A missing product path can terminate the single-instance process.

ZERO-PRICE VARIATION

A normal variation path can enter the cart with a zero base price.

PAGINATION RACES

Late responses can overwrite another category namespace.

COUPON / SHIPPING DRIFT

Multiple conflicting implementations disagree on money.

RUNTIME CONTRACT

Container runtime is older than a locked image dependency requires.

FALSE CONFIDENCE

Tests miss the financial trust boundary and depend on absent fixtures.

Security was only half the race

The winning separation came from ordinary-user correctness: zero-price variations, async pagination corruption, and exact framework control flow. A pure security sweep would have missed the benchmark-winning defects.

Quality and provider-native usage

Read this page carefully

Provider percentages are not tokens, dollars, or normalized compute. Only the two GLM runs share the same provider meter and are directly comparable.

GLM quality score



GLM session meter movement



GLM request count



Operator-observed, evidence pending. Full GLM used ~39% more requests for a 13.8-point quality advantage.

Astra

~2 pp

WEEKLY METER CONSUMED

99% left to 97% left; no visible 5-hour movement. Screenshot-backed operator observation.

Fable

~8%

CLAUDE 5-HOUR WINDOW

Operator-observed, evidence pending. Different provider and denominator.

Operational interpretation

Within the comparable GLM pair, the heavier run bought a materially better review. Across providers, no efficiency ranking is valid. Cross-provider percentages remain separate cards by design.

Why we trust the result

PINNED SPECIMEN

DockBloxx commit 059ccdba... certified by two independent methods; 362 files, zero symlinks.

MANIFEST AT EVERY GATE

362/362 SHA-256 entries verified after every run, during blind packaging, workspace build, and final compilation.

RAW ARTIFACT PRESERVATION

Every contestant report frozen and hashed before blind packaging. No post-hoc improvements.

SEALED RANDOM MAPPING

A/B/C/D drawn once from OS randomness; identities existed only in the sealed mapping.

PHYSICAL REFEREE ISOLATION

Sol worked outside the repo with no Git, identities, economics, run order, or mapping.

BLIND AUDIT BEFORE REVEAL

Jarvis verified every weighted score, grade, and rank before Tony opened the mapping.

What this benchmark supports - and does not

Supported

- Astra produced the strongest review on this frozen specimen.
- All four diagnosed the central architecture correctly.
- Validity, verification, and ordinary-user correctness decided the race.
- Referee source access materially changed adjudication.
- Target, raw reports, blind reports, and reveal chain stayed intact.

Not proven

- Not a universal model ranking.
- Not a cost or normalized efficiency result.
- Not a statistically reliable 5.6-point Astra/Fable margin.
- Not proof that conditional findings occur in production.
- Not standing Factory doctrine from one specimen.

Limitations carried forward

One repository, one run per model, Fable's prior exposure to an earlier DockBloxx review, Astra's deeper controlled-execution method, no live services, and heterogeneous provider meters.

The championship decision

On this production commerce specimen, Astra was the strongest reviewer. Fable was the strongest breadth partner. GLM 5.3 remained a capable senior pass. Flash found the architecture, but required more human verification.

Final ranking

RANK 1 Astra 95.6 A+	RANK 2 Fable 5.1 90.0 A	RANK 3 GLM 5.3 85.0 B	RANK 4 GLM Flash 71.2 C
---	--	--	--

Factory hypotheses - not doctrine

Principal verification

Astra for high-stakes claims where false certainty is expensive: payments, authorization, concurrency, incident triage.

Breadth + synthesis

Fable for broad system sweeps, security/financial trust boundaries, test critique, and actionable repair sequencing.

Routine review tiers

GLM 5.3 for strong core reviews; **Flash** for faster tactical passes only when a human verifies consequential claims.

README

Understand the result and the story in 60 seconds.

Executive PDF

Present the decision with charts, profiles, and integrity.

Full Markdown

Audit the entire experiment, evidence chain, and limitations.

Evidence beats assertion. Verification beats reputation.

One frozen codebase. Four review styles. One blind result.